



Road Accident Severity Classification

In [2]:

```
!pip install imbalanced-learn
```

```
Defaulting to user installation because normal site-packages is not writeable
```

```
Collecting imbalanced-learn
```

```
  Downloading imbalanced_learn-0.10.1-py3-none-any.whl (226 kB)
```

```
----- 226.0/226.0 kB 726.5 kB/s eta 0:0
```

```
0:00
```

```
Requirement already satisfied: scipy>=1.3.2 in c:\programdata\anaconda3\lib\site-packages (from imbalanced-learn) (1.9.1)
```

```
Requirement already satisfied: threadpoolctl>=2.0.0 in c:\programdata\anaconda3\lib\site-packages (from imbalanced-learn) (2.2.0)
```

```
Requirement already satisfied: numpy>=1.17.3 in c:\programdata\anaconda3\lib\site-packages (from imbalanced-learn) (1.21.5)
```

```
Requirement already satisfied: scikit-learn>=1.0.2 in c:\programdata\anaconda3\lib\site-packages (from imbalanced-learn) (1.0.2)
```

```
Collecting joblib>=1.1.1
```

```
  Downloading joblib-1.2.0-py3-none-any.whl (297 kB)
```

```
----- 298.0/298.0 kB 1.5 MB/s eta 0:0
```

```
0:00
```

```
Installing collected packages: joblib, imbalanced-learn
```

```
Successfully installed imbalanced-learn-0.10.1 joblib-1.2.0
```

In [3]:

```
#import the necessary Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import time
from collections import Counter
from imblearn.over_sampling import SMOTE
import matplotlib.ticker as ticker
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler,MinMaxScaler,LabelEncoder
from sklearn.pipeline import Pipeline
from sklearn.model_selection import RepeatedStratifiedKFold,GridSearchCV
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.ensemble import RandomForestClassifier,AdaBoostClassifier,GradientBoostingClassifier
from sklearn.metrics import accuracy_score,classification_report,confusion_matrix
from sklearn.model_selection import KFold # import KFold
import warnings
warnings.filterwarnings('ignore')
%matplotlib inline
```

In [4]:

```
df = pd.read_csv("RTA Dataset.csv")
```

Let's take a look at the dataset


In [5]:

```
df.head()
```

Out[5]:

	Time	Day_of_week	Age_band_of_driver	Sex_of_driver	Educational_level	Vehicle_driver_
0	17:02:00	Monday	18-30	Male	Above high school	Er
1	17:02:00	Monday	31-50	Male	Junior high school	Er
2	17:02:00	Monday	18-30	Male	Junior high school	Er
3	1:06:00	Sunday	18-30	Male	Junior high school	Er
4	1:06:00	Sunday	18-30	Male	Junior high school	Er

5 rows × 32 columns




In [6]:

```
df.sample(5)
```

Out[6]:

	Time	Day_of_week	Age_band_of_driver	Sex_of_driver	Educational_level	Vehicle_driv
2084	7:59:00	Friday	31-50	Male	Junior high school	
8376	20:00:00	Thursday	Over 51	Male	Elementary school	
8672	15:30:00	Friday	18-30	Male	Junior high school	
2825	1:05:00	Saturday	31-50	Male	Junior high school	
873	11:29:00	Saturday	Over 51	Male	Junior high school	

5 rows × 32 columns



In [7]:

```
df.shape
```

Out[7]:

```
(12316, 32)
```

In [8]:

```
df.columns
```

Out[8]:

```
Index(['Time', 'Day_of_week', 'Age_band_of_driver', 'Sex_of_driver',  
      'Educational_level', 'Vehicle_driver_relation', 'Driving_experience',  
      'Type_of_vehicle', 'Owner_of_vehicle', 'Service_year_of_vehicle',  
      'Defect_of_vehicle', 'Area_accident_occured', 'Lanes_or_Medians',  
      'Road_alignment', 'Types_of_Junction', 'Road_surface_type',  
      'Road_surface_conditions', 'Light_conditions', 'Weather_conditions',  
      'Type_of_collision', 'Number_of_vehicles_involved',  
      'Number_of_casualties', 'Vehicle_movement', 'Casualty_class',  
      'Sex_of_casualty', 'Age_band_of_casualty', 'Casualty_severity',  
      'Work_of_casualty', 'Fitness_of_casualty', 'Pedestrian_movement',  
      'Cause_of_accident', 'Accident_severity'],  
      dtype='object')
```


In [9]:

```
df.describe(include="all")
```

Out[9]:

	Time	Day_of_week	Age_band_of_driver	Sex_of_driver	Educational_level	Vehicle_d
count	12316	12316	12316	12316	11575	
unique	1074	7	5	3	7	
top	15:30:00	Friday	18-30	Male	Junior high school	
freq	120	2041	4271	11437	7619	
mean	NaN	NaN	NaN	NaN	NaN	
std	NaN	NaN	NaN	NaN	NaN	
min	NaN	NaN	NaN	NaN	NaN	
25%	NaN	NaN	NaN	NaN	NaN	
50%	NaN	NaN	NaN	NaN	NaN	
75%	NaN	NaN	NaN	NaN	NaN	
max	NaN	NaN	NaN	NaN	NaN	

11 rows × 32 columns



In [10]:

```
df.dtypes
```

Out[10]:

Time	object
Day_of_week	object
Age_band_of_driver	object
Sex_of_driver	object
Educational_level	object
Vehicle_driver_relation	object
Driving_experience	object
Type_of_vehicle	object
Owner_of_vehicle	object
Service_year_of_vehicle	object
Defect_of_vehicle	object
Area_accident_occured	object
Lanes_or_Medians	object
Road_allignment	object
Types_of_Junction	object
Road_surface_type	object
Road_surface_conditions	object
Light_conditions	object
Weather_conditions	object
Type_of_collision	object
Number_of_vehicles_involved	int64
Number_of_casualties	int64
Vehicle_movement	object
Casualty_class	object
Sex_of_casualty	object
Age_band_of_casualty	object
Casualty_severity	object
Work_of_casualty	object
Fitness_of_casualty	object
Pedestrian_movement	object
Cause_of_accident	object
Accident_severity	object
dtype:	object

In [11]:

df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 12316 entries, 0 to 12315
Data columns (total 32 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Time                                  12316 non-null  object
1   Day_of_week                          12316 non-null  object
2   Age_band_of_driver                  12316 non-null  object
3   Sex_of_driver                       12316 non-null  object
4   Educational_level                   11575 non-null  object
5   Vehicle_driver_relation             11737 non-null  object
6   Driving_experience                   11487 non-null  object
7   Type_of_vehicle                     11366 non-null  object
8   Owner_of_vehicle                    11834 non-null  object
9   Service_year_of_vehicle             8388 non-null   object
10  Defect_of_vehicle                   7889 non-null   object
11  Area_accident_occured               12077 non-null  object
12  Lanes_or_Medians                    11931 non-null  object
13  Road_allignment                     12174 non-null  object
14  Types_of_Junction                   11429 non-null  object
15  Road_surface_type                   12144 non-null  object
16  Road_surface_conditions              12316 non-null  object
17  Light_conditions                    12316 non-null  object
18  Weather_conditions                  12316 non-null  object
19  Type_of_collision                   12161 non-null  object
20  Number_of_vehicles_involved          12316 non-null  int64
21  Number_of_casualties                 12316 non-null  int64
22  Vehicle_movement                     12008 non-null  object
23  Casualty_class                       12316 non-null  object
24  Sex_of_casualty                     12316 non-null  object
25  Age_band_of_casualty                 12316 non-null  object
26  Casualty_severity                    12316 non-null  object
27  Work_of_casualty                     9118 non-null   object
28  Fitness_of_casualty                  9681 non-null   object
29  Pedestrian_movement                 12316 non-null  object
30  Cause_of_accident                   12316 non-null  object
31  Accident_severity                   12316 non-null  object
dtypes: int64(2), object(30)
memory usage: 3.0+ MB
```

In [12]:

```
# convert the 'Date' column to datetime format
df['Time'] = pd.to_datetime(df['Time'])
```

In [13]:

```
df.duplicated()
```

Out[13]:

```
0      False
1      False
2      False
3      False
4      False
...
12311   False
12312   False
12313   False
12314   False
12315   False
Length: 12316, dtype: bool
```

In [14]:

```
df.duplicated().sum()
```

Out[14]:

```
0
```

In [15]:

```
df.groupby('Accident_severity').size()
```

Out[15]:

```
Accident_severity
Fatal injury      158
Serious Injury   1743
Slight Injury    10415
dtype: int64
```

Data Preprocessing

In [16]:

```
df.isnull().sum()
```

Out[16]:

Time	0
Day_of_week	0
Age_band_of_driver	0
Sex_of_driver	0
Educational_level	741
Vehicle_driver_relation	579
Driving_experience	829
Type_of_vehicle	950
Owner_of_vehicle	482
Service_year_of_vehicle	3928
Defect_of_vehicle	4427
Area_accident_occured	239
Lanes_or_Medians	385
Road_allignment	142
Types_of_Junction	887
Road_surface_type	172
Road_surface_conditions	0
Light_conditions	0
Weather_conditions	0
Type_of_collision	155
Number_of_vehicles_involved	0
Number_of_casualties	0
Vehicle_movement	308
Casualty_class	0
Sex_of_casualty	0
Age_band_of_casualty	0
Casualty_severity	0
Work_of_casualty	3198
Fitness_of_casualty	2635
Pedestrian_movement	0
Cause_of_accident	0
Accident_severity	0

dtype: int64

In []:


We can summarize the table as:

Number of observations: 12316 Number of columns: 32 Memory Usage: 3.0+ MB Number of int columns: 2
Number of object columns: 30 Number of columns with missing values: 16

Numerical data analysis

In [17]:

```
df.hist(layout=(1,6), figsize=(30,5))  
plt.show()
```



In [18]:


```
df['Number_of_casualties'].value_counts()
```

Out[18]:

```
1    8397  
2    2290  
3     909  
4     394  
5     207  
6      89  
7      22  
7       8  
Name: Number_of_casualties, dtype: int64
```


In [19]:

```
plt.figure(figsize=(10,7))  
sns.boxplot(data=df, y='Number_of_vehicles_involved', x='Number_of_casualties')  
plt.show()
```




In [20]:

```
sns.boxplot(data=df, y='Number_of_casualties')  
plt.show()
```



In [21]:

```
sns.boxplot(data=df, y='Number_of_vehicles_involved')  
plt.show()
```



In [22]:


```
df['Number_of_vehicles_involved']
```

Out[22]:

```
0      2
1      2
2      2
3      2
4      2
..
12311   2
12312   2
12313   1
12314   2
12315   2
Name: Number_of_vehicles_involved, Length: 12316, dtype: int64
```


In [23]:

```
sns.scatterplot(x=df['Number_of_vehicles_involved'], y=df['Number_of_casualties'])
plt.show()
```




In [24]:

```
sns.pairplot(df[['Number_of_vehicles_involved', 'Number_of_casualties']])  
plt.show()
```



In [25]:

```
correlation_matrix = df[['Number_of_vehicles_involved', 'Number_of_casualties']].corr()  
sns.heatmap(correlation_matrix, annot=True)  
plt.show()
```




In this heatmap, we can find that these are not much extremely correlated variables

Categorical data analysis


In [26]:

```
plt.figure(figsize=(10,7))  
plt.pie(x=df['Accident_severity'].value_counts().values,  
        labels=df['Accident_severity'].value_counts().index,  
        autopct='%2.2f%%')  
plt.show()
```



In [27]:

```
# creating a facet grid with columns as survived=0 and survived=1
grid = sns.FacetGrid(data=df, col='Accident_severity', height=4, aspect=1, sharey=False)
# mapping bar plot and the data on to the grid
grid.map(sns.countplot, 'Number_of_vehicles_involved', palette=['black', 'brown', 'orange'])
plt.show()
```



In [28]:

```
df.columns
```

Out[28]:

```
Index(['Time', 'Day_of_week', 'Age_band_of_driver', 'Sex_of_driver',
      'Educational_level', 'Vehicle_driver_relation', 'Driving_experience',
      'Type_of_vehicle', 'Owner_of_vehicle', 'Service_year_of_vehicle',
      'Defect_of_vehicle', 'Area_accident_occured', 'Lanes_or_Medians',
      'Road_alignment', 'Types_of_Junction', 'Road_surface_type',
      'Road_surface_conditions', 'Light_conditions', 'Weather_conditions',
      'Type_of_collision', 'Number_of_vehicles_involved',
      'Number_of_casualties', 'Vehicle_movement', 'Casualty_class',
      'Sex_of_casualty', 'Age_band_of_casualty', 'Casualty_severity',
      'Work_of_casualty', 'Fitness_of_casualty', 'Pedestrian_movement',
      'Cause_of_accident', 'Accident_severity'],
      dtype='object')
```

In [29]:

```
# dropping columns that can cause imbalance while imputation
lists=['Vehicle_driver_relation', 'Work_of_casualty', 'Fitness_of_casualty', 'Day_of_week']
df.drop(columns = lists, inplace=True)
```

In [30]:

```
df.shape
```

Out[30]:

```
(12316, 19)
```

In [31]:

```
df.columns
```

Out[31]:

```
Index(['Age_band_of_driver', 'Driving_experience', 'Type_of_vehicle',  
      'Area_accident_occured', 'Lanes_or_Medians', 'Road_allignment',  
      'Types_of_Junction', 'Road_surface_conditions', 'Light_conditions',  
      'Weather_conditions', 'Type_of_collision',  
      'Number_of_vehicles_involved', 'Number_of_casualties',  
      'Vehicle_movement', 'Casualty_class', 'Age_band_of_casualty',  
      'Pedestrian_movement', 'Cause_of_accident', 'Accident_severity'],  
      dtype='object')
```

Filling missing values

In [32]:

```
# fill missing values with mean column values  
df['Driving_experience'].fillna(df['Driving_experience'].mode()[0], inplace=True)  
df['Age_band_of_driver'].fillna(df['Age_band_of_driver'].mode()[0], inplace=True)  
df['Type_of_vehicle'].fillna(df['Type_of_vehicle'].mode()[0], inplace=True)  
df['Area_accident_occured'].fillna(df['Area_accident_occured'].mode()[0], inplace=True)  
df['Road_allignment'].fillna(df['Road_allignment'].mode()[0], inplace=True)  
df['Type_of_collision'].fillna(df['Type_of_collision'].mode()[0], inplace=True)  
df['Vehicle_movement'].fillna(df['Vehicle_movement'].mode()[0], inplace=True)  
df['Lanes_or_Medians'].fillna(df['Lanes_or_Medians'].mode()[0], inplace=True)  
df['Types_of_Junction'].fillna(df['Types_of_Junction'].mode()[0], inplace=True)
```

In [33]:

```
df.isnull().sum()
```

Out[33]:

Age_band_of_driver	0
Driving_experience	0
Type_of_vehicle	0
Area_accident_occured	0
Lanes_or_Medians	0
Road_allignment	0
Types_of_Junction	0
Road_surface_conditions	0
Light_conditions	0
Weather_conditions	0
Type_of_collision	0
Number_of_vehicles_involved	0
Number_of_casualties	0
Vehicle_movement	0
Casualty_class	0
Age_band_of_casualty	0
Pedestrian_movement	0
Cause_of_accident	0
Accident_severity	0
dtype:	int64

In [34]:

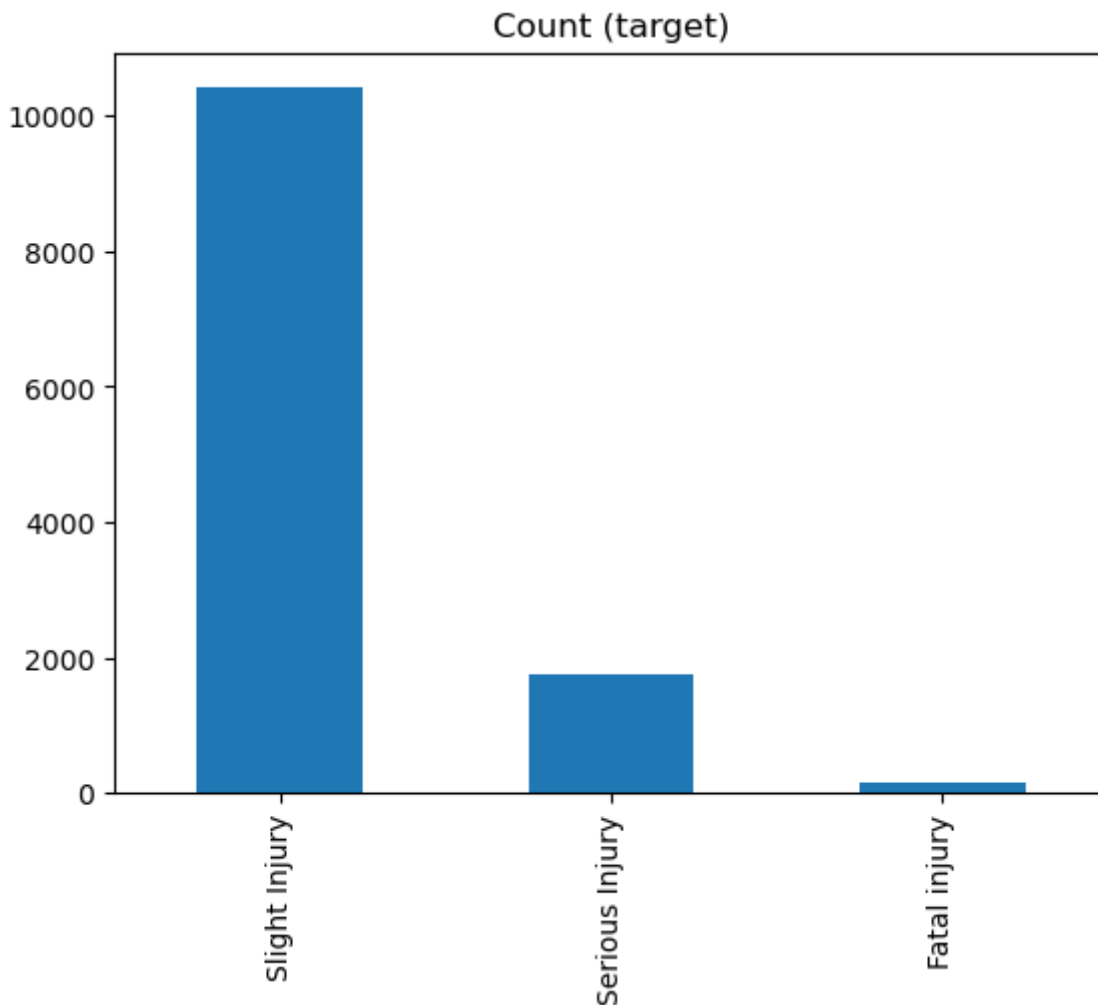
```
target_count = df['Accident_severity'].value_counts()
print('Class 0:', target_count[0])
print('Class 1:', target_count[1])
print('Proportion:', round(target_count[0] / target_count[1], 2), ': 1')

target_count.plot(kind='bar', title='Count (target)');
```

Class 0: 10415

Class 1: 1743

Proportion: 5.98 : 1



Encoding

In [35]:


```
from sklearn.preprocessing import LabelEncoder          #or one hot encoder
LE = LabelEncoder()
df=df.apply(LE.fit_transform)                          #categorical values to integers
```

In [36]:

```
plt.figure(figsize=[25,15])
sns.heatmap(df.corr(),annot=True)
```

Out[36]:

<AxesSubplot:>



In [37]:

```
for col in df.select_dtypes(include='object'):
    if df[col].nunique() <= 22:
        sns.countplot(y=col, data=df)
        plt.show()
```

Upsampling

In [38]:

```
x = df.drop('Accident_severity', axis=1)
y = df['Accident_severity']

xtrain, xtest, ytrain, ytest = train_test_split(x, y, test_size=0.3, random_state=42)
print(xtrain.shape, xtest.shape, ytrain.shape, ytest.shape)
```

```
(8621, 18) (3695, 18) (8621,) (3695,)
```

In [39]:

```
# upsampling using smote

counter = Counter(ytrain)

print("=====")

for k,v in counter.items():
    per = 100*v/len(ytrain)
    print(f"Class= {k}, n={v} ({per:.2f}%)")

oversample = SMOTE()
xtrain, ytrain = oversample.fit_resample(xtrain, ytrain)

counter = Counter(ytrain)

print("=====")

for k,v in counter.items():
    per = 100*v/len(ytrain)
    print(f"Class= {k}, n={v} ({per:.2f}%)")

print("=====")

print("Upsampled data shape: ", xtrain.shape, ytrain.shape)
```

```
=====
Class= 2, n=7324 (84.96%)
Class= 1, n=1191 (13.82%)
Class= 0, n=106 (1.23%)
=====
Class= 2, n=7324 (33.33%)
Class= 1, n=7324 (33.33%)
Class= 0, n=7324 (33.33%)
=====
Upsampled data shape: (21972, 18) (21972,)
```

Splitting test and train data

In [40]:

```
x=df.drop(columns=["Accident_severity"])
y=df["Accident_severity"]
```

In [41]:

```
models={"LogisticRegression":LogisticRegression(),
        "DecisionTreeClassifier":DecisionTreeClassifier(),
        "SVM":SVC(),
        "KNeighborsClassifier":KNeighborsClassifier(),
        "GNB":GaussianNB(),
        "RandomForestClassifier":RandomForestClassifier(),
        "AdaBoostClassifier":AdaBoostClassifier(),
        "GradientBoostingClassifier":GradientBoostingClassifier(),
        }
```

In [42]:

```
# models,x,y,scaleFlag=0,1,2
def modelAccuracy(models,x,y,scaleFlag):
    #train/Test
    xtrain,xtest,ytrain,ytest=train_test_split(x,y,test_size=0.2,random_state=0)
    acc_result={}
    for name,model in models.items():
        #pipeline
        #1.Transformer -> 2.Model
        if(scaleFlag==1):
            model_pipeline=Pipeline([('MinMaxScaler',MinMaxScaler()),('model',model)])
        elif(scaleFlag==2):
            model_pipeline=Pipeline([('StandardScaler',StandardScaler()),('model',model)])
        else:
            model_pipeline=Pipeline([('model',model)])
        #training/testing on model pipeline
        model_fit=model_pipeline.fit(xtrain,ytrain)
        ypred=model_fit.predict(xtest)
        acc=accuracy_score(ytest,ypred)
        print("The Accuracy for ",name," is :",acc)
        acc_result[name]=acc
    return acc_result
```

In [43]:

```
def bestModel(result):
    high=0
    for name,acc in result.items():
        if acc>high:
            high=acc
            model_name=name
    print("Best Model is ",model_name," with accuaracy =>",high)
```

In [44]:

```
def bestParams(model,param,xtrain,ytrain):
    #cv
    cv=RepeatedStratifiedKFold(n_splits=5,n_repeats=3)
    grid_cv=GridSearchCV(estimator=model,param_grid=param,cv=cv,scoring="f1_weighted")
    res=grid_cv.fit(xtrain,ytrain)
    print("Best Parameters are ",res.best_params_)
    print("Best Accuracy is ",res.best_score_)
```

In [45]:

bestParams

Out[45]:

<function __main__.bestParams(model, param, xtrain, ytrain)>

In [46]:

```
acc=modelAccuracy(models,x,y,1)
```

```
The Accuracy for LogisticRegression is : 0.84375
The Accuracy for DecisionTreeClassifier is : 0.7329545454545454
The Accuracy for SVM is : 0.84375
The Accuracy for KNeighborsClassifier is : 0.8262987012987013
The Accuracy for GNB is : 0.8145292207792207
The Accuracy for RandomForestClassifier is : 0.8482142857142857
The Accuracy for AdaBoostClassifier is : 0.8425324675324676
The Accuracy for GradientBoostingClassifier is : 0.8486201298701299
```

In [47]:

```
bestModel(acc)
```

```
Best Model is GradientBoostingClassifier with accuaracy => 0.8486201298701299
```

In [48]:

```
model=RandomForestClassifier()
params={"n_estimators" : [100,200],
        "criterion" : ["gini","entropy"]}
bestParams(model,params,xtrain,ytrain)
```

```
Best Parameters are {'criterion': 'entropy', 'n_estimators': 200}
Best Accuracy is 0.9194693388169922
```

In [49]:

```
#retrain the model with best parameters
model=RandomForestClassifier(criterion="entropy",n_estimators=200)
model.fit(xtrain,ytrain)
ypred=model.predict(xtest)
```


In [50]:


```
#Final Evaluation
print(accuracy_score(ytest,ypred))
print(classification_report(ytest,ypred))
cm=confusion_matrix(ytest,ypred)
sns.heatmap(cm,annot=True)
```

0.7691474966170501

	precision	recall	f1-score	support
0	0.03	0.02	0.02	52
1	0.24	0.20	0.22	552
2	0.85	0.88	0.87	3091
accuracy			0.77	3695
macro avg	0.38	0.37	0.37	3695
weighted avg	0.75	0.77	0.76	3695

Out[50]:

<AxesSubplot:>



In []:

In []: